

# Spotify Hit Predictor

Christopher El-Khouri

625.740: Data Mining

November 17, 2021

## 1. Introduction

Spotify is one of the most popular music streaming services in the world with 381 million active monthly users as of October 2021 (Spotify Technology S.A, 2021). In addition to being a digital music streaming platform, Spotify created a platform called Spotify for Developers (Spotify AB, 2021) that provides users with a set of tools to retrieve Spotify measured audio features for the different tracks on the application. Audio features such as danceability, energy, instrumentalness, tempo, and more can be extracted.

In this project, we are going to develop a classification model that uses Spotify audio feature data to predict whether or not a track was a hit i.e. achieved mainstream popularity. The anticipated challenges are as follows:

- Different audio features may have different effects on popularity depending on the time of release of a musical track. For example: The 1980s may have had a higher mainstream preference for more danceable music compared to the 1970s so a “dancey” track released in the 70s may have not been a hit but would have been if it were released in the 1980s.
- Different audio features may have different effects on popularity depending on the genre of a musical track. For example: A rock song with a high danceability measure may not be a hit as much as a Hip hop song with the same danceability.
- A track could be a hit because of the artist and album it was released with, regardless of the audio features and this may skew our analysis and modeling.

The performance of 4 algorithms will be assessed on the same dataset; Logistic Regression, Decision Trees, K-Nearest Neighbors, and Feedforward Neural Networks.

## 2. Background

This report continues my previous research conducted in the Spring of 2020 titled “Music Taste Across Decades: Spotify Audio Feature Analysis of 33,354 Musical Pieces” (El Khouri, 2020). The research paper was for the Data Visualization course (605.662) and it performs a statistical and visual analysis of audio feature data of 33,354 musical tracks whose audio features have been measured by Spotify. Hypothesis testing was performed testing the relationships between genres, popularity, audio features, and time. The audio features; instrumentalness, danceability, valence, and energy were found to be correlated with popularity. Based on the results, I expect to build a reliable model that is able to predict the hit status of an arbitrary track using its audio features, decade, and genre.

### 3. Algorithms and Experimental Methods

The sections below elaborate upon the algorithms and experimental methods used in this project.

#### 3.1. Logistic Regression

In the field of machine learning, linear models are the simplest of algorithms and function by generating a line of best fit to successfully predict the value of the target variable using the feature attributes. A typical linear model has the following form:

$$(1) \quad f(x) = w_0 + \sum_{i=1}^d w_i x_i = w_0 + \mathbf{w}_i^T \mathbf{x}$$

Logistic Regression is a supervised linear machine learning algorithm that uses the logistic distribution to classify a dependent variable. In equation (2),  $P(C_i|x)$  refers to the probability that input feature data  $x$  belongs to a class  $C_i$ .

Taking  $f(x)$  from (1):

$$(2) \quad P(C_i|x) = \frac{\exp(f(x_i))}{\sum_{j=1}^K \exp(f(x_j))}, i = 1, \dots, K$$

The Logistic Regression algorithm used in this project is the statsmodels Logit algorithm (Seabold & Perktold 2010).

#### Backward Elimination

The Logistic Regression model is initially run with all the features from the dataset. Based on the resulting p-values of the individual features, backward elimination is implemented where the feature with the highest p-value greater than 0.05 is eliminated. The algorithm is run on the modified dataset again until there are no more features with p-values greater than 0.05. The final modified dataset is the dataset that will be used to assess the performance of the rest of the algorithms.

#### 3.2. Decision Trees

Decision Trees are a type of supervised learning algorithms that can be used for classification and regression. When used for classification, these trees have a target variable that takes discrete values. The branches are observations about the attributes and the leaves are the different class labels that our data could hold.

An example of a classification tree is shown Figure 1 (Colton, 2006), the branches are based on the following attributes: Parents Visiting, Weather, and Money. The leaves are based on the following class labels: Cinema, Play Tennis, Stay in, and Shopping. The decision tree is used by starting from the top, "Parents Visiting" in this case, and traverses until we arrive at a leaf.

The Decision Tree algorithm used in this project is the scikit-learn tree algorithm (Pedregosa et al., 2011).

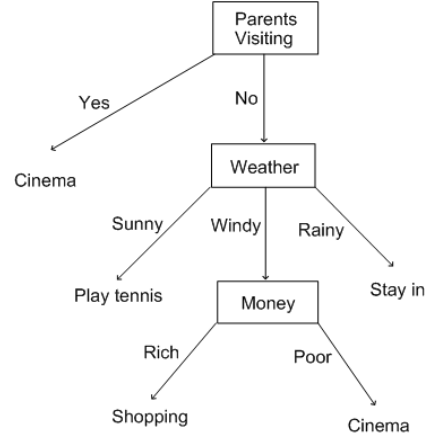


Figure 1: Classification Tree Example

### Minimal Cost-Complexity Pruning

Decision Trees can be pruned to avoid overfitting and improve model performance. The scikit-learn algorithm implements Minimal Cost-Complexity Pruning (Pedregosa et al., 2011) and is demonstrated by equation (3) below:

$$(3) \quad R_{\alpha}(T) = R(T) + \alpha |\tilde{T}|$$

$\alpha$ : Complexity Parameter

$|\tilde{T}|$ : Number of leaves in tree  $T$

$R(T)$ : Total misclassification rate of the leaves

Minimal Cost-Complexity Pruning finds the subtree that minimizes  $R_{\alpha}(T)$ . 10% of our data is extracted to form a Tuning Set. The decision tree algorithm is implemented on our Tuning Set with  $\alpha$  varying from 0 to 0.5 to fine tune the complexity parameter and maximize model performance. The  $\alpha$  corresponding to the maximal performance on the Tuning Set is implemented for the remaining 90% of the data.

### 3.3. K-Nearest Neighbors

K-Nearest Neighbors (K-NN) is a non-parametric method initially proposed by Thomas Cover (1967) that can be used for classification or regression. When used for classification, the algorithm takes the input data and returns the most common class among its k nearest neighbors from the training data. The figure below (Yadnesh, 2020) shows the working principle of K-NN.

The K-NN classification algorithm used in this project is the scikit-learn KNeighborsClassifier (Pedregosa et al., 2011). 10% of our data is extracted to form a Tuning Set. The K-NN algorithm is implemented on our Tuning Set with varying values of k to fine tune the number of neighbors and maximize model performance. The k corresponding to the maximal performance on the Tuning Set is implemented for the remaining 90% of the data.




Figure 2: K-NN Example

### 3.4. Feedforward Neural Networks

Feedforward Neural Networks are a type of Artificial Neural Networks (ANN) that generally utilize perceptrons, which are supervised linear learning algorithms that can be used for classification or regression.

Figure 3 below shows a simple perceptron:




Figure 3: Simple Perceptron

The output  $y$  has the form of a typical linear model:

$$(4) \quad y = w_0 + \sum_{i=1}^d w_i x_i$$

Multilayer Perceptrons (MLP) are perceptrons with 1 or more hidden layers, where the neurons in one layer are connected to the neurons in the other in a feedforward type structure. Figure 4 shows a 1-layer feedforward neural network with multiple outputs,  $z_1 \dots z_h$  are the activation functions that typically take sigmoid, tanh, or gaussian transformations.




Figure 4: 1-layered MLP

The Multilayer Perceptron classification algorithm used in this project is the scikit-learn MLPClassifier (Pedregosa et al., 2011). 10% of our data is extracted to form a Tuning Set. The MLP algorithm is implemented on our Tuning Set with varying numbers of hidden dimensions to maximize model performance. The hidden dimensions corresponding to the maximal performance on the Tuning Set is implemented for the remaining 90% of the data.

### 3.5. 5-fold cross-validation

For each of the models, the data underwent 5-fold cross-validation, where the data is split into five parts (or folds) and the model is trained on each of the folds accordingly. In each run, one of the five folds is used as a test set and the remaining four folds are concatenated and used as the training set. The performance of each model is determined by calculating the classification errors of each of the folds by dividing the number of incorrect predictions over the total size of the folds, then calculating the mean of the errors.

## 4. Dataset

The dataset used is the same data set used in my previous research paper Music Taste Across Decades: Spotify Audio Feature Analysis of 33,354 Musical Pieces (El Khouri, 2020), which is based on The Spotify Hit Predictor Dataset (1960-2019) (Ansari, 2020). In my research paper, I preprocessed the dataset to group the tracks by genre and removed outliers accordingly. Further details to the preprocessing can be found on the paper.

The dataset used in this project initially contained 33,354 tracks, 31 features, and 2 classes: Hit (denoted by '1') and Flop (denoted by '0'). Further preprocessing was performed in order to prepare our data for modeling which resulted in the final dataset containing 58 features. The feature attributes; loudness, tempo, duration\_ms, chorus\_hit, and sections were scaled to have their values between 0 and 1. Dummy variables were created for the features time\_signature, genres, Decade, key\_full. The features of the final dataset are listed below.

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	danceability	33354 non-null	float64
1	energy	33354 non-null	float64
2	loudness	33354 non-null	float64
3	speechiness	33354 non-null	float64
4	acousticness	33354 non-null	float64
5	instrumentalness	33354 non-null	float64
6	liveness	33354 non-null	float64
7	valence	33354 non-null	float64
8	tempo	33354 non-null	float64
9	duration_ms	33354 non-null	float64
10	chorus_hit	33354 non-null	float64
11	sections	33354 non-null	float64
12	target	33354 non-null	int64
13	time_signature_1	33354 non-null	uint8
14	time_signature_3	33354 non-null	uint8
15	time_signature_4	33354 non-null	uint8
16	time_signature_5	33354 non-null	uint8
17	genres_Caribbean	33354 non-null	uint8
18	genres_Classical	33354 non-null	uint8
19	genres_Country	33354 non-null	uint8
20	genres_Easy listening	33354 non-null	uint8
21	genres_Electronic	33354 non-null	uint8
22	genres_Folk	33354 non-null	uint8
23	genres_Hip hop	33354 non-null	uint8
24	genres_Jazz	33354 non-null	uint8
25	genres_Latin	33354 non-null	uint8
26	genres_Metal	33354 non-null	uint8
27	genres_Pop	33354 non-null	uint8
28	genres_R&B	33354 non-null	uint8
29	genres_Rock	33354 non-null	uint8
30	Decade_10s	33354 non-null	uint8
31	Decade_60s	33354 non-null	uint8
32	Decade_70s	33354 non-null	uint8
33	Decade_80s	33354 non-null	uint8
34	Decade_90s	33354 non-null	uint8
35	key_full_A Minor	33354 non-null	uint8
36	key_full_B Major	33354 non-null	uint8
37	key_full_B Minor	33354 non-null	uint8
38	key_full_Bb Major	33354 non-null	uint8
39	key_full_Bb Minor	33354 non-null	uint8
40	key_full_C Major	33354 non-null	uint8
41	key_full_C Minor	33354 non-null	uint8
42	key_full_C# Major	33354 non-null	uint8
43	key_full_C# Minor	33354 non-null	uint8
44	key_full_D Major	33354 non-null	uint8
45	key_full_D Minor	33354 non-null	uint8
46	key_full_E Major	33354 non-null	uint8
47	key_full_E Minor	33354 non-null	uint8
48	key_full_Eb Major	33354 non-null	uint8
49	key_full_Eb Minor	33354 non-null	uint8
50	key_full_F Major	33354 non-null	uint8
51	key_full_F Minor	33354 non-null	uint8
52	key_full_F# Major	33354 non-null	uint8
53	key_full_F# Minor	33354 non-null	uint8
54	key_full_G Major	33354 non-null	uint8
55	key_full_G Minor	33354 non-null	uint8
56	key_full_G# Major	33354 non-null	uint8
57	key_full_G# Minor	33354 non-null	uint8

However, after performing backward elimination, the dataset that will proceed to be used with the remaining models contains the following features:

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	danceability	33354 non-null	float64
1	energy	33354 non-null	float64
2	loudness	33354 non-null	float64
3	speechiness	33354 non-null	float64
4	acousticness	33354 non-null	float64
5	instrumentalness	33354 non-null	float64
6	liveness	33354 non-null	float64
7	valence	33354 non-null	float64
8	tempo	33354 non-null	float64
9	chorus_hit	33354 non-null	float64
10	genres_Caribbean	33354 non-null	uint8
11	genres_Classical	33354 non-null	uint8
12	genres_Country	33354 non-null	uint8
13	genres_Easy listening	33354 non-null	uint8
14	genres_Electronic	33354 non-null	uint8
15	genres_Folk	33354 non-null	uint8
16	genres_Hip hop	33354 non-null	uint8
17	genres_Latin	33354 non-null	uint8
18	genres_Metal	33354 non-null	uint8
19	genres_Pop	33354 non-null	uint8
20	genres_R&B	33354 non-null	uint8
21	genres_Rock	33354 non-null	uint8
22	Decade_10s	33354 non-null	uint8
23	Decade_60s	33354 non-null	uint8
24	Decade_70s	33354 non-null	uint8
25	Decade_80s	33354 non-null	uint8
26	Decade_90s	33354 non-null	uint8
27	key_full_A Minor	33354 non-null	uint8
28	key_full_B Minor	33354 non-null	uint8
29	key_full_Bb Major	33354 non-null	uint8
30	key_full_C Major	33354 non-null	uint8
31	key_full_C# Major	33354 non-null	uint8
32	key_full_C# Minor	33354 non-null	uint8
33	key_full_D Minor	33354 non-null	uint8
34	key_full_E Major	33354 non-null	uint8
35	key_full_E Minor	33354 non-null	uint8
36	key_full_Eb Major	33354 non-null	uint8
37	key_full_F Major	33354 non-null	uint8
38	key_full_F# Major	33354 non-null	uint8
39	key_full_G# Major	33354 non-null	uint8
40	target	33354 non-null	int64

A total of 40 features including the target variable. Further details can be found in the supporting documents.

## 5. Results

The results from our different models are shown on the table below:

Table 1: Results Table

Model	Measure	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean
Logistic Regression	Error	17.3%	17.0%	17.8%	16.6%	17.3%	<b>17.2%</b>
	False Positives	22.9%	23.2%	23.5%	22.5%	23.4%	<b>23.1%</b>
	False Negatives	11.5%	10.6%	11.8%	10.7%	11.0%	<b>11.1%</b>
Logistic Regression with Backward Elimination	Error	17.4%	16.9%	18.0%	16.5%	17.4%	<b>17.2%</b>
	False Positives	22.9%	23.2%	23.7%	22.3%	23.8%	<b>23.2%</b>
	False Negatives	11.8%	10.4%	12.0%	10.7%	10.9%	<b>11.2%</b>
Decision Tree (alpha=0.0001)	Error	18.6%	18.9%	17.8%	19.0%	19.7%	<b>18.8%</b>
	False Positives	20.5%	20.2%	19.5%	19.4%	21.9%	<b>20.3%</b>
	False Negatives	16.7%	17.6%	16.1%	18.6%	17.4%	<b>17.2%</b>
K-NN (n=5)	Error	17.7%	18.5%	17.5%	17.4%	18.8%	<b>18.0%</b>
	False Positives	23.8%	25.1%	23.1%	23.5%	24.1%	<b>23.9%</b>
	False Negatives	11.7%	11.7%	11.6%	11.1%	13.3%	<b>11.9%</b>
MLP (Hidden Layer Sizes=[3,4])	Error	14.8%	15.3%	15.8%	14.5%	16.3%	<b>15.3%</b>
	False Positives	17.5%	19.0%	19.1%	19.4%	18.0%	<b>18.6%</b>
	False Negatives	12.1%	11.4%	12.4%	9.4%	14.6%	<b>12.0%</b>

The results can be summarized as follows:

Table 2: Summary of Results

Model	Measure	Mean
Logistic Regression	Error	17.2%
	False Positives	23.1%
	False Negatives	11.1%
Logistic Regression with Backward Elimination	Error	17.2%
	False Positives	23.2%
	False Negatives	11.2%
Decision Tree (alpha=0.0001)	Error	18.8%
	False Positives	20.3%
	False Negatives	17.2%
K-NN (n=5)	Error	18.0%
	False Positives	23.9%
	False Negatives	11.9%
MLP (Hidden Layer Sizes=[3,4])	Error	<b>15.3%</b>
	False Positives	<b>18.6%</b>
	False Negatives	<b>12.0%</b>



## 6. Discussion

Since the goal was to predict the hit status of a track, the measures of error and false positives/negatives are important. Having a low false positive rate is more important than a low false negative rate because incorrectly predicting that a track is a hit rather than incorrectly predicting that a track is not a hit may have costly implications.

Referring to Table 2, the MLP algorithm gave us the strongest model for this data, with an error rate of 15.3 %, this model has an accuracy of approximately 85% and a false positive rate of only 18.6%. This means that the model will correctly predict a hit ~80% of the time. Since the hit status of a musical track is a sort of psychological/social science phenomenon, as it is a matter of subjective opinion, an accuracy of 85% is outstanding.

Although K-NN and Logistic Regression had slightly superior error rates over Decision Trees, Decision Trees had superior false positive rates. The slight raise in error rate in return for a superior false positive rate make Decision Trees the next best model for this case in my opinion.

## 7. Conclusion

In conclusion, I believe that the goal of this project was achieved with a model accuracy of 85%. Referring to our anticipated challenges mentioned in the introduction of this report, I believe that the first two challenges were accounted for by incorporating genre and decade into our models. However, the third was not, neither ‘artist’ nor ‘album’ was a feature incorporated into our models. Moreover, a track may have the technical audio features required to be a ‘hit’ but the artist or record label may not have the marketing strategy necessary for it to achieve that.

Further modelling analysis can be performed incorporating artist, album, and possibly record label. Furthermore, additional models can be assessed in an effort to develop more accurate predictive models.

Regardless, the results so far are significant enough to encourage further developments in the understanding of herd music preference with respect to scientifically measured audio data. This can lead to implications regarding the marketing of music from a platform, record label, and a musician’s perspective.

## 8. References

Spotify Technology S.A.

[Spotify Q2 Press Release 2021](#), page 1.  
2021.

Spotify AB.

[Spotify for Developers](#).  
2021.

Christopher El Khouri.

[Music Taste Across Decades: Spotify Audio Feature Analysis of 33,354 Musical Pieces](#).  
2020.

Skipper Seabold and Josef Perktold.

[statsmodels: Econometric and statistical modeling with python](#).  
In *Proceedings of the 9th Python in Science Conference*,  
pages 92-96, 2010.

Simon Colton.

Lecture 6: Decision Tree Learning.

Artificial Intelligence in Computational Art and Design, School of Interactive Arts  
and Technology.

Simon Fraser University, British Columbia, Canada, 2006.

Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand  
Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent  
Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu  
Brucher, Matthieu Perrot, and Édouard Duchesnay.

[Scikit-learn: Machine Learning in Python](#).

*Journal of Machine Learning Research*, 12(85):2825–2830, 2011.

Thomas Cover and Peter Hart.

Nearest neighbor pattern classification

IEEE Transactions on Information Theory, 13(1):21-27, 1967.

Yadnesh.

[k-Nearest Neighbors in Machine Learning \(k-NN\)](#).

2020.

Farooq Ansari.

[The Spotify Hit Predictor Dataset \(1960-2019\)](#).

2020.